

LSTM Model for playing Hangman

Aditya Raut

September 11th, 2024

1 Task description

The provided file 'words_250000_train.txt' contains a total of 227,300 words made of small alphabet letters. The Jupyter notebook `Hangman_run_experiments.ipynb` contains all the necessary functions to play Hangman games after loading a pre-trained model. The task is to maximize the win rate playing the game with 6 lives, i.e., you can make 6 incorrect letter guesses.

2 Proposed strategy

I trained a bidirectional LSTM neural network created using PyTorch locally on a custom dataset created from the given word dictionary. The best model as per my judgment was used for predictions.

- **Dataset creation -**

- For every **word** in given dictionary, I created 15 data points of 'guessed made in the past'.
- Every generated set of guesses has the following realistic game conditions -
 - * at least 1 letter remaining to guess correctly in the **word**
 - * at most 5 incorrectly guessed letters (at most 5 lives lost)
- Encode the 'masked **word**' of length n containing `--` symbols as a PyTorch tensor of size $n \times 27$, where each of the n rows has 1 at `ord(letter)-97` or at index 26 for `--` and 0 everywhere else.
- The ideal output of the model should be equal probability for all remaining letters in the **word**.
- Due to creating 15 different guesses per word, we have over 3.4 million data points.

- **Model architecture -**

- I used a bidirectional LSTM that first takes input of 'masked **word**' encoded as above.
 - * LSTM has `num_layers = 3`, `hidden_size = 256`, `dropout = 0.2`
- This LSTM gives output of size 512, which is $2 * \text{hidden_size}$.
- We add a dropout layer of 0.4 before connecting this fully to a linear layer.
- We then concatenate a 1-0 indicator tensor of size 26 indicating all guesses made.
- This tensor of size 538 is passed to a linear layer of dimensions (538,128).
- Another dropout layer with 0.4 value between fully connected layers.
- Finally a linear layer of size (128,26), which gives probabilities for all letters.

- **Loss criterion - CrossEntropyLoss**, since this uses a **softmax** function to compare with labels.

- **Training parameters -**

- Random 9:1 ratio data split for training and validation, `batch_size = 500`.
- ADAM Optimizer with initial `learning_rate=1e-3`, `weight_decay=1e-5`
- StepLR scheduler that cuts learning rate by factor of $\frac{1}{2}$ every 10 epochs.
- Trained for a total of 250 epochs.

- **Final model and guess selection -**

- Model at epoch 241 had the least validation loss, and was selected for final submissions.
- The highest probability un-guessed letter from the model is returned by `guess(self, word)`.

3 Important files

1. `Hangman_models_training.ipynb` - Dataset creation and model training
2. `Hangman_run_experiments.ipynb` - to load pre-trained models and run on random or chosen words
3. `Hangman_Strategy_Aditya_Raut.pdf` - This PDF, description of strategy

Very large sized outputs in the training Jupyter notebook are cleared for better readability.